

AYUSH KARNEWAR

AI/ML Engineer Intern | Generative AI | LLM & NLP Applications | RAG Systems

+91 9960670950 | ayushkarnewar369@gmail.com | [linkedin.com/in/ayush-karnewar](https://www.linkedin.com/in/ayush-karnewar) | github.com/ayush-karnewar | www.ayushkarnewar.in

PROFESSIONAL SUMMARY

AI/ML Engineer with hands-on experience developing and deploying machine learning models and generative AI applications. Skilled in Natural Language Processing (NLP) pipelines, Large Language Models (LLMs), Retrieval-Augmented Generation (RAG), and prompt engineering. Proficient in Python, FastAPI, and Docker for building scalable AI systems. Experienced with AWS, Azure, and Google Cloud (Vertex AI). Strong focus on model performance optimization, backend development, and system design.

CORE COMPETENCIES

Machine Learning · Deep Learning · Generative AI · Large Language Models (LLMs) · NLP Pipelines · Retrieval-Augmented Generation (RAG) · Prompt Engineering · Python · FastAPI · TensorFlow · PyTorch · Scikit-learn · Docker · AWS · Azure · Google Cloud · API Development · Backend Development · System Design

EXPERIENCE

AI/ML Intern — AI Deployment & Automation

Mar 2026

AI Deployment & Automation Internship

- Boosted model response relevance by ~35% and cut inference latency by ~20% by building and deploying RAG pipelines on AWS, Azure, and Google Cloud.
- Shipped 5+ production-ready AI applications end-to-end using FastAPI and Docker, reducing deployment time through containerized pipelines.

Lead Coordinator — International Conference on Recent Trends in Engineering & Technology

Feb 2026

ICRTET

- Managed 1000+ research paper submissions across 6 months; oversaw reviewer allocation, author communication, and end-to-end conference execution for 200+ participants.

AI/ML Intern — AWS Cloud Virtual Internship

Dec 2024

Amazon Web Services (AWS)

- Raised LLM response accuracy by ~25% by engineering optimized prompt workflows and fine-tuning generative AI pipelines on AWS (EC2, S3).
- Accelerated 3+ real-world NLP delivery cycles by integrating third-party APIs for text generation and automated summarization.

PROJECTS

IDiot — AI-Powered IoT Circuit Simulator

Apr 2026

Stack: Python, JavaScript, LLMs, Generative AI

Cut IoT circuit design time by ~50% by building a canvas-based simulator with drag-and-drop components, AI layout suggestions, and LLM-powered Arduino code generation.

Vision AI — Smart Photo Gallery

Mar 2026

Stack: Python, NLP, Deep Learning, Face Recognition

Reduced manual photo sorting effort by ~70% across 1000+ images by building an AI gallery using face recognition and NLP-based semantic search with automated clustering.

ResearchHub AI — Research Discovery Platform

Jan 2026

Stack: Python, RAG, LLMs, NLP, FastAPI

Slashed literature review time by ~40% by building a RAG-powered research assistant with NLP semantic search and automated academic paper summarization.

DUSTIN — Voice-Enabled Mental Health Companion

Dec 2025

Stack: Python, RAG, LLMs, Speech Recognition, FastAPI

Delivered 24/7 accessible mental health support by engineering a voice-enabled RAG assistant with NLP models for emotionally context-aware conversations.

TECHNICAL SKILLS

AI / ML:	Machine Learning, Deep Learning, NLP, LLMs, RAG, Generative AI, Prompt Engineering, Model Fine-tuning
Frameworks:	FastAPI, TensorFlow, PyTorch, Scikit-learn, LangChain
Cloud & DevOps:	AWS (EC2, S3), Azure, Google Cloud / Vertex AI, Docker, CI/CD, Git/GitHub
Languages & DB:	Python, JavaScript, C, C++ MySQL, MongoDB, Supabase

EDUCATION

B.Tech — CSE (AI-ML), D.Y. Patil College of Engineering & Technology, Kolhapur | 2023–2027 | CGPA: 8.5/10

ADDITIONAL

Certifications: Google AI (2025), AWS Cloud (2025), Cisco Data Science (2025) · **Research:** EMG-based AAC System; Persona.exe; The Gravity Hoax · **Languages:** Hindi, Marathi, English, Spanish